

Local Linear Estimation

We want to estimate the gradient of a function f , whose gradient is supposed sparse. We suppose that we are able to gather evaluations $Y = f(X)$ of the function of interest.

Local Linear Estimation of the Gradient

If we locally approximate the function by its Taylor expansion, we can estimate the gradient by solving

$$(\tilde{m}_k(x), \tilde{\beta}_k(x)) \in \arg \min_{(m, \beta) \in \mathbb{R}^{D+1}} \sum_{i: X_i \in \text{KNN}(x)} (Y_i - m - \beta^\top (X_i - x))^2 + \lambda \|\beta\|_1$$

Theorem 1.

Let $n \geq 1$ and $k \geq 1$ such that $\bar{\tau}_k \leq \tau_0$. Let $\delta \in (0, 1)$ and set $\lambda = \bar{\tau}_k(\sqrt{2\sigma^2 \log(16D/\delta)/k} + L_2 \bar{\tau}_k^2)$. Then, we have with probability larger than $1 - \delta$,

$$\|\tilde{\beta}_k(x) - \beta(x)\|_2 \leq (24)^2 \sqrt{\#S_x} \left(\bar{\tau}_k^{-1} \sqrt{\frac{2\sigma^2 \log(16D/\delta)}{k}} + L_2 \bar{\tau}_k \right),$$

as soon as $C_1 \#S_x \log(Dn/\delta) \leq k \leq C_2 n$, $\bar{\tau}_k^2 \leq (b_f^2 / (C_3 \#S_x L^2) \wedge \tau_0^2)$, where C_1 , C_2 and C_3 are universal constants.

Our bounds make use of the *sparsity* of ∇f and only require a neighbourhood consisting of the K nearest points.

The use of a KNN neighbourhood gives us:

- Robustness to the data
- Ease of calibration
- Possibility to reuse past computations

Variable Selection

Require: (X, Y) : training set, Node: indexes of points in the node

- 1: $\nabla m(X_i) \leftarrow$ estimated gradient at X_i , $\forall i \in \text{Node}$ using (1)
- 2: $\omega \leftarrow \sum_{i \in \text{Node}} |\nabla m(X_i)|$
- 3: $K \leftarrow$ sample \sqrt{D} dimensions in $\{1, \dots, d\}$ with probability weights $\propto \omega$
- 4: $k, c \leftarrow$ best threshold c and dimension k
- 5: **return** k, c

Algorithm 1: Node Splitting for Gradient Guided Trees

Variable Selection

Dataset	Description		Loss	
	n	D	RF	GGF
Wisconsin	569	30	0.0352	0.0345
Heart	303	13	0.128	0.124
Diamonds	53940	23	680033	664265
Gasoline	60	401	0.678	0.512
SDSS	10000	8	$0.872 \cdot 10^{-3}$	$0.776 \cdot 10^{-3}$

Table 1: Performance of the two random forest algorithms on a 50-folds cross validation.

Gradient Free Optimization

Require: x_0 : initial guess, f : function $\mathbb{R}^D \rightarrow \mathbb{R}$, M : budget

- 1: $X \leftarrow X_1, \dots, X_M$ with $X_i \sim \mathcal{N}(x_0, \varepsilon \times I_D)$
- 2: $Y \leftarrow f(X) := f(X_1), \dots, f(X_M)$
- 3: **while** not StoppingCondition **do**
- 4: $m, \Delta \leftarrow$ estimated gradient at x w.r.t X, Y using (1)
- 5: $X \leftarrow X, X_1, \dots, X_M$ with $X_i \sim \mathcal{N}(\text{GradientStep}(x, \Delta), \varepsilon \times I_D)$
- 6: $Y \leftarrow f(X)$
- 7: $x \leftarrow \arg \min_{X_i} \{f(X_i)\}$
- 8: **return** x

Algorithm 2: Estimated Gradient Descent

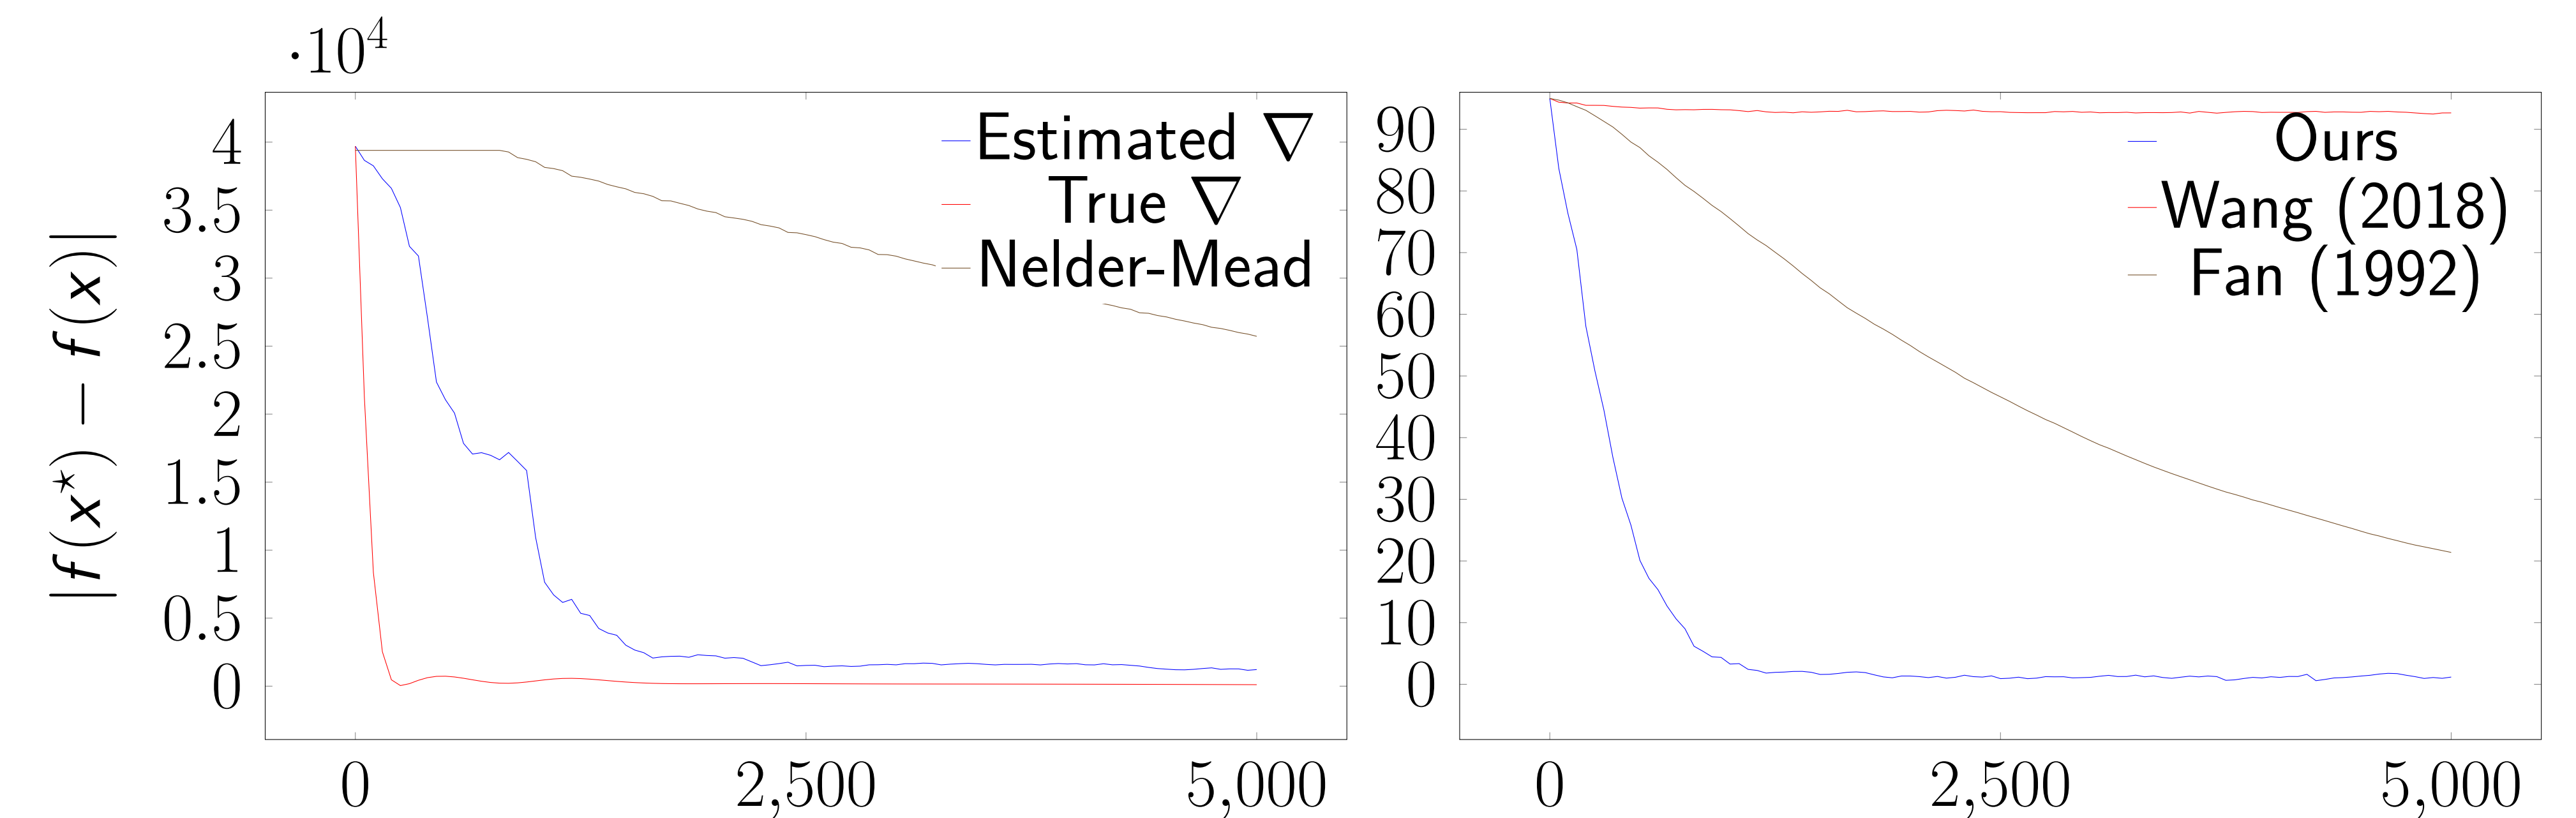


Figure 1: Gradient Descent on the sparse noisy Rosenbrock function for $d = 100$.

